

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import matplotlib as mpl
%matplotlib inline
mpl.style.use('ggplot')
```

```
In [2]: car=pd.read_csv('quikr_car.csv')
```

```
In [3]: car.head()
```

Out[3]:

	name	company	year	Price	kms_driven	fuel_type
0	Hyundai Santro Xing XO eRLX Euro III	Hyundai	2007	80,000	45,000 kms	Petrol
1	Mahindra Jeep CL550 MDI	Mahindra	2006	4,25,000	40 kms	Diesel
2	Maruti Suzuki Alto 800 Vxi	Maruti	2018	Ask For Price	22,000 kms	Petrol
3	Hyundai Grand i10 Magna 1.2 Kappa VTVT	Hyundai	2014	3,25,000	28,000 kms	Petrol
4	Ford EcoSport Titanium 1.5L TDCi	Ford	2014	5,75,000	36,000 kms	Diesel

```
In [4]: car.shape
```

Out[4]: (892, 6)

```
In [5]: car.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 892 entries, 0 to 891
Data columns (total 6 columns):
#   Column          Non-Null Count  Dtype
---  -
0   name            892 non-null   object
1   company         892 non-null   object
2   year            892 non-null   object
3   Price           892 non-null   object
4   kms_driven      840 non-null   object
5   fuel_type       837 non-null   object
dtypes: object(6)
memory usage: 41.9+ KB
```

Creating backup copy

```
In [6]: backup=car.copy()
```

Quality

- names are pretty inconsistent
- names have company names attached to it

- some names are spam like 'Maruti Ertiga showroom condition with' and 'Well mentained Tata Sumo'
- company: many of the names are not of any company like 'Used', 'URJENT', and so on.
- year has many non-year values
- year is in object. Change to integer
- Price has Ask for Price
- Price has commas in its prices and is in object
- kms_driven has object values with kms at last.
- It has nan values and two rows have 'Petrol' in them
- fuel_type has nan values

Cleaning Data

year has many non-year values

```
In [7]: car=car[car['year'].str.isnumeric()]
```

year is in object. Change to integer

```
In [8]: car['year']=car['year'].astype(int)
```

Price has Ask for Price

```
In [9]: car=car[car['Price']!='Ask For Price']
```

Price has commas in its prices and is in object

```
In [10]: car['Price']=car['Price'].str.replace(',', '').astype(int)
```

kms_driven has object values with kms at last.

```
In [11]: car['kms_driven']=car['kms_driven'].str.split().str.get(0).str.replace(' ', '')
```

It has nan values and two rows have 'Petrol' in them

```
In [12]: car=car[car['kms_driven'].str.isnumeric()]
```

```
In [13]: car['kms_driven']=car['kms_driven'].astype(int)
```

fuel_type has nan values

```
In [14]: car=car[~car['fuel_type'].isna()]
```

```
In [15]: car.shape
```

```
Out[15]: (816, 6)
```

name and company had spammed data...but with the previous cleaning, those rows got removed.

Company does not need any cleaning now. Changing car names. Keeping only the first three words

```
In [16]: car['name']=car['name'].str.split().str.slice(start=0,stop=3).str.join(' ')
```

Resetting the index of the final cleaned data

```
In [17]: car=car.reset_index(drop=True)
```

Cleaned Data

```
In [18]: car
```

```
Out[18]:
```

	name	company	year	Price	kms_driven	fuel_type
0	Hyundai Santro Xing	Hyundai	2007	80000	45000	Petrol
1	Mahindra Jeep CL550	Mahindra	2006	425000	40	Diesel
2	Hyundai Grand i10	Hyundai	2014	325000	28000	Petrol
3	Ford EcoSport Titanium	Ford	2014	575000	36000	Diesel
4	Ford Figo	Ford	2012	175000	41000	Diesel
...
811	Maruti Suzuki Ritz	Maruti	2011	270000	50000	Petrol
812	Tata Indica V2	Tata	2009	110000	30000	Diesel
813	Toyota Corolla Altis	Toyota	2009	300000	132000	Petrol
814	Tata Zest XM	Tata	2018	260000	27000	Diesel
815	Mahindra Quanto C8	Mahindra	2013	390000	40000	Diesel

816 rows × 6 columns

```
In [19]: car.to_csv('Cleaned_Car_data.csv')
```

In [20]: `car.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 816 entries, 0 to 815
Data columns (total 6 columns):
#   Column      Non-Null Count  Dtype
---  -
0   name        816 non-null    object
1   company     816 non-null    object
2   year        816 non-null    int32
3   Price       816 non-null    int32
4   kms_driven  816 non-null    int32
5   fuel_type   816 non-null    object
dtypes: int32(3), object(3)
memory usage: 28.8+ KB
```

In [21]: `car.describe(include='all')`

Out[21]:

	name	company	year	Price	kms_driven	fuel_type
count	816	816	816.000000	8.160000e+02	816.000000	816
unique	254	25	NaN	NaN	NaN	3
top	Maruti Suzuki Swift	Maruti	NaN	NaN	NaN	Petrol
freq	51	221	NaN	NaN	NaN	428
mean	NaN	NaN	2012.444853	4.117176e+05	46275.531863	NaN
std	NaN	NaN	4.002992	4.751844e+05	34297.428044	NaN
min	NaN	NaN	1995.000000	3.000000e+04	0.000000	NaN
25%	NaN	NaN	2010.000000	1.750000e+05	27000.000000	NaN
50%	NaN	NaN	2013.000000	2.999990e+05	41000.000000	NaN
75%	NaN	NaN	2015.000000	4.912500e+05	56818.500000	NaN
max	NaN	NaN	2019.000000	8.500003e+06	400000.000000	NaN

In []:

In [22]: `car=car[car['Price']<6000000]`


Checking relationship of Company with Price

In [23]: `car['company'].unique()`

Out[23]: array(['Hyundai', 'Mahindra', 'Ford', 'Maruti', 'Skoda', 'Audi', 'Toyota', 'Renault', 'Honda', 'Datsun', 'Mitsubishi', 'Tata', 'Volkswagen', 'Chevrolet', 'Mini', 'BMW', 'Nissan', 'Hindustan', 'Fiat', 'Force', 'Mercedes', 'Land', 'Jaguar', 'Jeep', 'Volvo'], dtype=object)

In [24]: `import seaborn as sns`

```
In [25]: plt.subplots(figsize=(15,7))  
ax=sns.boxplot(x='company',y='Price',data=car)  
ax.set_xticklabels(ax.get_xticklabels(),rotation=40,ha='right')  
plt.show()
```



Checking relationship of Year with Price

```
In [26]: plt.subplots(figsize=(20,10))
ax=sns.swarmplot(x='year',y='Price',data=car)
ax.set_xticklabels(ax.get_xticklabels(),rotation=40,ha='right')
plt.show()
```

C:\Users\Sayaib Sarkar\AppData\Local\Programs\Python\Python310\lib\site-packages\seaborn\categorical.py:1296: UserWarning: 9.3% of the points cannot be placed; you may want to decrease the size of the markers or use stripplot.

warnings.warn(msg, UserWarning)

C:\Users\Sayaib Sarkar\AppData\Local\Programs\Python\Python310\lib\site-packages\seaborn\categorical.py:1296: UserWarning: 6.8% of the points cannot be placed; you may want to decrease the size of the markers or use stripplot.

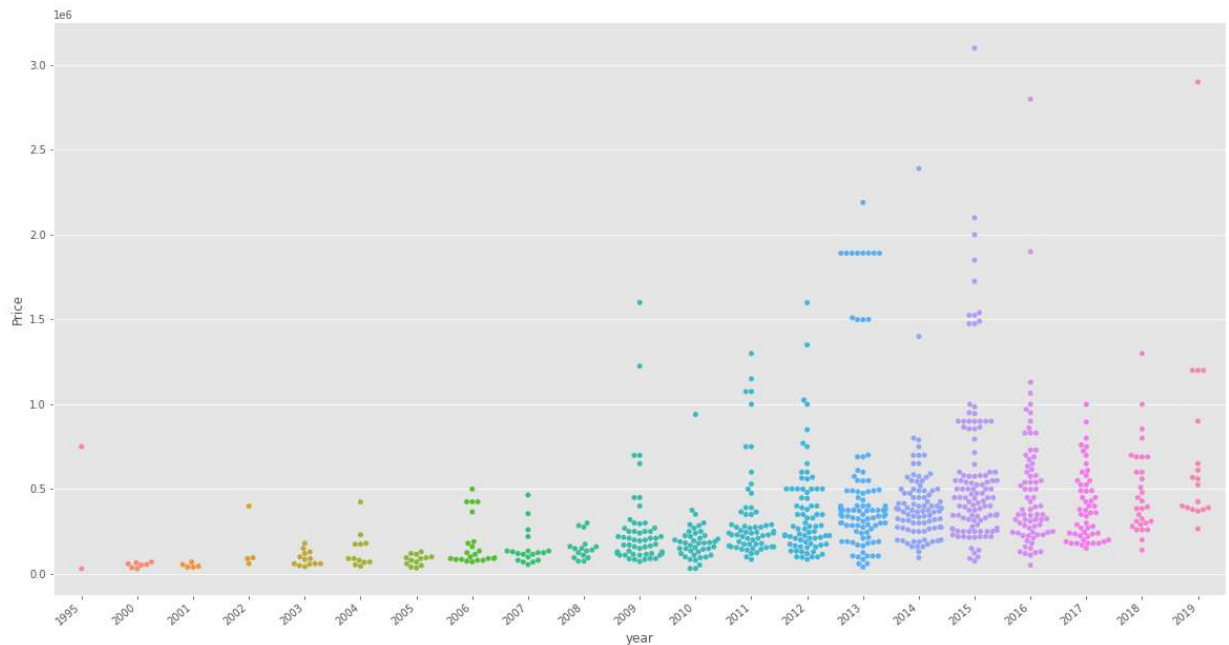
warnings.warn(msg, UserWarning)

C:\Users\Sayaib Sarkar\AppData\Local\Programs\Python\Python310\lib\site-packages\seaborn\categorical.py:1296: UserWarning: 10.6% of the points cannot be placed; you may want to decrease the size of the markers or use stripplot.

warnings.warn(msg, UserWarning)

C:\Users\Sayaib Sarkar\AppData\Local\Programs\Python\Python310\lib\site-packages\seaborn\categorical.py:1296: UserWarning: 5.5% of the points cannot be placed; you may want to decrease the size of the markers or use stripplot.


warnings.warn(msg, UserWarning)



Checking relationship of kms_driven with Price

```
In [27]: sns.relplot(x='kms_driven',y='Price',data=car,height=7,aspect=1.5)
```


```
Out[27]: <seaborn.axisgrid.FacetGrid at 0x15d7632beb0>
```



Checking relationship of Fuel Type with Price

```
In [28]: plt.subplots(figsize=(14,7))  
sns.boxplot(x='fuel_type',y='Price',data=car)
```


```
Out[28]: <AxesSubplot:xlabel='fuel_type', ylabel='Price'>
```



Relationship of Price with FuelType, Year and Company mixed


```
In [29]: ax=sns.relplot(x='company',y='Price',data=car,hue='fuel_type',size='year',height=
ax.set_xticklabels(rotation=40,ha='right')
```

```
Out[29]: <seaborn.axisgrid.FacetGrid at 0x15d7632b550>
```



Extracting Training Data

```
In [30]: X=car[['name','company','year','kms_driven','fuel_type']]
y=car['Price']
```

```
In [31]: X
```

```
Out[31]:
```

	name	company	year	kms_driven	fuel_type
0	Hyundai Santro Xing	Hyundai	2007	45000	Petrol
1	Mahindra Jeep CL550	Mahindra	2006	40	Diesel
2	Hyundai Grand i10	Hyundai	2014	28000	Petrol
3	Ford EcoSport Titanium	Ford	2014	36000	Diesel
4	Ford Figo	Ford	2012	41000	Diesel
...
811	Maruti Suzuki Ritz	Maruti	2011	50000	Petrol
812	Tata Indica V2	Tata	2009	30000	Diesel
813	Toyota Corolla Altis	Toyota	2009	132000	Petrol
814	Tata Zest XM	Tata	2018	27000	Diesel
815	Mahindra Quanto C8	Mahindra	2013	40000	Diesel

815 rows × 5 columns

```
In [32]: y.shape
```

```
Out[32]: (815,)
```

Applying Train Test Split

```
In [33]: from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.2)
```

```
In [34]: from sklearn.linear_model import LinearRegression
```

```
In [35]: from sklearn.preprocessing import OneHotEncoder
from sklearn.compose import make_column_transformer
from sklearn.pipeline import make_pipeline
from sklearn.metrics import r2_score
```

Creating an OneHotEncoder object to contain all the possible categories

```
In [36]: ohe=OneHotEncoder()
ohe.fit(X[['name','company','fuel_type']])
```

```
Out[36]: OneHotEncoder()
```

Creating a column transformer to transform categorical columns

```
In [37]: column_trans=make_column_transformer((OneHotEncoder(categories=ohe.categories_),
                                                remainder='passthrough'))
```

Linear Regression Model

```
In [38]: lr=LinearRegression()
```

Making a pipeline

```
In [39]: pipe=make_pipeline(column_trans,lr)
```

Fitting the model

```
In [40]: pipe.fit(X_train,y_train)
```

```
Out[40]: Pipeline(steps=[('columntransformer',
                           ColumnTransformer(remainder='passthrough',
                                                transformers=[('onehotencoder',
                                                                OneHotEncoder(categories=[array(['Audi A3 Cabriolet', 'Audi A4 1.8', 'Audi A4 2.0', 'Audi A6 2.0',
                                                                'Audi A8', 'Audi Q3 2.0', 'Audi Q5 2.0', 'Audi Q7', 'BMW 3 Series',
                                                                'BMW 5 Series', 'BMW 7 Series', 'BMW X1', 'BMW X1 sDrive20d',
                                                                'BMW X1 xDrive20d', 'Chevrolet Beat', 'Chevrolet Beat...
                                                                array(['Audi', 'BMW', 'Chevrolet', 'Datsun', 'Fiat', 'Force', 'Ford',
                                                                'Hindustan', 'Honda', 'Hyundai', 'Jaguar', 'Jeep', 'Land',
                                                                'Mahindra', 'Maruti', 'Mercedes', 'Mini', 'Mitsubishi', 'Nissan',
                                                                'Renault', 'Skoda', 'Tata', 'Toyota', 'Volkswagen', 'Volvo'],
                                                                dtype=object)),
                                                                array(['Diesel', 'LPG', 'Petrol'], dtype=object))]),
                           ('name', 'company',
                                'fuel_type'])]),
               ('linearregression', LinearRegression())])
```

```
In [41]: y_pred=pipe.predict(X_test)
```

Checking R2 Score

```
In [42]: r2_score(y_test,y_pred)
```

```
Out[42]: 0.786133164572156
```

Finding the model with a random state of TrainTestSplit where the model was found to give almost 0.92 as r2_score

```
In [43]: scores=[]
         for i in range(1000):
             X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.1,random_state=i)
             lr=LinearRegression()
             pipe=make_pipeline(column_trans,lr)
             pipe.fit(X_train,y_train)
             y_pred=pipe.predict(X_test)
             scores.append(r2_score(y_test,y_pred))
```

```
In [44]: np.argmax(scores)
```

```
Out[44]: 655
```

```
In [45]: scores[np.argmax(scores)]
```

```
Out[45]: 0.920087093218515
```

```
In [46]: pipe.predict(pd.DataFrame(columns=X_test.columns,data=np.array(['Maruti Suzuki Sw  
Out[46]: array([400642.51767152])
```

The best model is found at a certain random state

```
In [47]: X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.1,random_state=np.  
lr=LinearRegression()  
pipe=make_pipeline(column_trans,lr)  
pipe.fit(X_train,y_train)  
y_pred=pipe.predict(X_test)  
r2_score(y_test,y_pred)
```

```
Out[47]: 0.920087093218515
```

```
In [48]: import pickle
```

```
In [49]: pickle.dump(pipe,open('LinearRegressionModel.pkl','wb'))
```

```
In [50]: pipe.predict(pd.DataFrame(columns=['name','company','year','kms_driven','fuel_type  
Out[50]: array([416107.38610967])
```

```
In [51]: pipe.steps[0][1].transformers[0][1].categories[0]
```

```
Out[51]: array(['Audi A3 Cabriolet', 'Audi A4 1.8', 'Audi A4 2.0', 'Audi A6 2.0',
'Audi A8', 'Audi Q3 2.0', 'Audi Q5 2.0', 'Audi Q7', 'BMW 3 Series',
'BMW 5 Series', 'BMW 7 Series', 'BMW X1', 'BMW X1 sDrive20d',
'BMW X1 xDrive20d', 'Chevrolet Beat', 'Chevrolet Beat Diesel',
'Chevrolet Beat LS', 'Chevrolet Beat LT', 'Chevrolet Beat PS',
'Chevrolet Cruze LTZ', 'Chevrolet Enjoy', 'Chevrolet Enjoy 1.4',
'Chevrolet Sail 1.2', 'Chevrolet Sail UVA', 'Chevrolet Spark',
'Chevrolet Spark 1.0', 'Chevrolet Spark LS', 'Chevrolet Spark LT',
'Chevrolet Tavera LS', 'Chevrolet Tavera Neo', 'Datsun GO T',
'Datsun Go Plus', 'Datsun Redi GO', 'Fiat Linea Emotion',
'Fiat Petra ELX', 'Fiat Punto Emotion', 'Force Motors Force',
'Force Motors One', 'Ford EcoSport', 'Ford EcoSport Ambiente',
'Ford EcoSport Titanium', 'Ford EcoSport Trend',
'Ford Endeavor 4x4', 'Ford Fiesta', 'Ford Fiesta SXi', 'Ford Figo',
'Ford Figo Diesel', 'Ford Figo Duratorq', 'Ford Figo Petrol',
'Ford Fusion 1.4', 'Ford Ikon 1.3', 'Ford Ikon 1.6',
'Hindustan Motors Ambassador', 'Honda Accord', 'Honda Amaze',
'Honda Amaze 1.2', 'Honda Amaze 1.5', 'Honda Brio', 'Honda Brio V',
'Honda Brio VX', 'Honda City', 'Honda City 1.5', 'Honda City SV',
'Honda City VX', 'Honda City ZX', 'Honda Jazz S', 'Honda Jazz VX',
'Honda Mobilio', 'Honda Mobilio S', 'Honda WR V', 'Hyundai Accent',
'Hyundai Accent Executive', 'Hyundai Accent GLE',
'Hyundai Accent GLX', 'Hyundai Creta', 'Hyundai Creta 1.6',
'Hyundai Elantra 1.8', 'Hyundai Elantra SX', 'Hyundai Elite i20',
'Hyundai Eon', 'Hyundai Eon D', 'Hyundai Eon Era',
'Hyundai Eon Magna', 'Hyundai Eon Sportz', 'Hyundai Fluidic Verna',
'Hyundai Getz', 'Hyundai Getz GLE', 'Hyundai Getz Prime',
'Hyundai Grand i10', 'Hyundai Santro', 'Hyundai Santro AE',
'Hyundai Santro Xing', 'Hyundai Sonata Transform', 'Hyundai Verna',
'Hyundai Verna 1.4', 'Hyundai Verna 1.6', 'Hyundai Verna Fluidic',
'Hyundai Verna Transform', 'Hyundai Verna VGT',
'Hyundai Xcent Base', 'Hyundai Xcent SX', 'Hyundai i10',
'Hyundai i10 Era', 'Hyundai i10 Magna', 'Hyundai i10 Sportz',
'Hyundai i20', 'Hyundai i20 Active', 'Hyundai i20 Asta',
'Hyundai i20 Magna', 'Hyundai i20 Select', 'Hyundai i20 Sportz',
'Jaguar XE XE', 'Jaguar XF 2.2', 'Jeep Wrangler Unlimited',
'Land Rover Freelander', 'Mahindra Bolero DI',
'Mahindra Bolero Power', 'Mahindra Bolero SLE',
'Mahindra Jeep CL550', 'Mahindra Jeep MM', 'Mahindra KUV100',
'Mahindra KUV100 K8', 'Mahindra Logan', 'Mahindra Logan Diesel',
'Mahindra Quanto C4', 'Mahindra Quanto C8', 'Mahindra Scorpio',
'Mahindra Scorpio 2.6', 'Mahindra Scorpio LX',
'Mahindra Scorpio S10', 'Mahindra Scorpio S4',
'Mahindra Scorpio SLE', 'Mahindra Scorpio SLX',
'Mahindra Scorpio VLX', 'Mahindra Scorpio Vlx',
'Mahindra Scorpio W', 'Mahindra TUV300 T4', 'Mahindra TUV300 T8',
'Mahindra Thar CRDe', 'Mahindra XUV500', 'Mahindra XUV500 W10',
'Mahindra XUV500 W6', 'Mahindra XUV500 W8', 'Mahindra Xylo D2',
'Mahindra Xylo E4', 'Mahindra Xylo E8', 'Maruti Suzuki 800',
'Maruti Suzuki A', 'Maruti Suzuki Alto', 'Maruti Suzuki Baleno',
'Maruti Suzuki Celerio', 'Maruti Suzuki Ciaz',
'Maruti Suzuki Dzire', 'Maruti Suzuki Eeco',
'Maruti Suzuki Ertiga', 'Maruti Suzuki Esteem',
'Maruti Suzuki Estilo', 'Maruti Suzuki Maruti',
'Maruti Suzuki Omni', 'Maruti Suzuki Ritz', 'Maruti Suzuki S',
```

```
'Maruti Suzuki SX4', 'Maruti Suzuki Stingray',  
'Maruti Suzuki Swift', 'Maruti Suzuki Versa',  
'Maruti Suzuki Vitara', 'Maruti Suzuki Wagon', 'Maruti Suzuki Zen',  
'Mercedes Benz A', 'Mercedes Benz B', 'Mercedes Benz C',  
'Mercedes Benz GLA', 'Mini Cooper S', 'Mitsubishi Lancer 1.8',  
'Mitsubishi Pajero Sport', 'Nissan Micra XL', 'Nissan Micra XV',  
'Nissan Sunny', 'Nissan Sunny XL', 'Nissan Terrano XL',  
'Nissan X Trail', 'Renault Duster', 'Renault Duster 110',  
'Renault Duster 110PS', 'Renault Duster 85', 'Renault Duster 85PS',  
'Renault Duster RxL', 'Renault Kwid', 'Renault Kwid 1.0',  
'Renault Kwid RXT', 'Renault Lodgy 85', 'Renault Scala RxL',  
'Skoda Fabia', 'Skoda Fabia 1.2L', 'Skoda Fabia Classic',  
'Skoda Laura', 'Skoda Octavia Classic', 'Skoda Rapid Elegance',  
'Skoda Superb 1.8', 'Skoda Yeti Ambition', 'Tata Aria Pleasure',  
'Tata Bolt XM', 'Tata Indica', 'Tata Indica V2', 'Tata Indica eV2',  
'Tata Indigo CS', 'Tata Indigo LS', 'Tata Indigo LX',  
'Tata Indigo Marina', 'Tata Indigo eCS', 'Tata Manza',  
'Tata Manza Aqua', 'Tata Manza Aura', 'Tata Manza ELAN',  
'Tata Nano', 'Tata Nano Cx', 'Tata Nano GenX', 'Tata Nano LX',  
'Tata Nano Lx', 'Tata Sumo Gold', 'Tata Sumo Grande',  
'Tata Sumo Victa', 'Tata Tiago Revotorq', 'Tata Tiago Revotron',  
'Tata Tigor Revotron', 'Tata Venture EX', 'Tata Vista Quadrajet',  
'Tata Zest Quadrajet', 'Tata Zest XE', 'Tata Zest XM',  
'Toyota Corolla', 'Toyota Corolla Altis', 'Toyota Corolla H2',  
'Toyota Etios', 'Toyota Etios G', 'Toyota Etios GD',  
'Toyota Etios Liva', 'Toyota Fortuner', 'Toyota Fortuner 3.0',  
'Toyota Innova 2.0', 'Toyota Innova 2.5', 'Toyota Qualis',  
'Volkswagen Jetta Comfortline', 'Volkswagen Jetta Highline',  
'Volkswagen Passat Diesel', 'Volkswagen Polo',  
'Volkswagen Polo Comfortline', 'Volkswagen Polo Highline',  
'Volkswagen Polo Highline1.2L', 'Volkswagen Polo Trendline',  
'Volkswagen Vento Comfortline', 'Volkswagen Vento Highline',  
'Volkswagen Vento Konekt', 'Volvo S80 Summum'], dtype=object)
```

In []: